Digital Image Processing Assignment 1: Image Reading, Enhancement and Resampling

Ian Tsai 313553058
Department of Computer Science
National Yang Ming Chiao Tung University
Hsinchu, Taiwan
Email: iantsai.cs13@nycu.edu.tw

Abstract—This report presents the implementation and results of three fundamental digital image processing techniques: image reading, image enhancement through point operations, and image resampling using interpolation methods. The implementation successfully reads RAW and BMP format images, applies log transform, gamma correction, and negative transform for enhancement, and performs image resizing using nearest neighbor and bilinear interpolation. Experimental results demonstrate the effectiveness of each technique with visual comparisons and quantitative analysis.

Index Terms—Digital Image Processing, Point Operations, Image Enhancement, Interpolation, Resampling

Source Code: https://github.com/IaTsai/MMIP-HW1.git

I. Introduction

Digital image processing is a fundamental field in computer vision and multimedia applications. This assignment focuses on implementing basic image processing operations that form the foundation for more complex algorithms. The complete implementation is publicly available at https://github.com/IaTsai/MMIP-HW1.git. The objectives include understanding different image formats, implementing enhancement techniques through point operations, and realizing image resampling methods.

The implementation covers three main areas:

- Image reading from RAW and standard formats
- Point-based enhancement operations
- Image resampling with different interpolation methods

II. Methodology

A. Image Reading

The image reading module handles two types of formats:

1) RAW Format: RAW images contain uncompressed pixel data in row-major order. The implementation reads binary data and reshapes it into a 512×512 grayscale matrix:

$$I[i, j] = \text{raw\_data}[i \times width + j] \tag{1}$$

2) BMP Format: BMP images are read using the PIL library, converting them to grayscale if necessary. The centered 10×10 pixel values are extracted for verification purposes.

B. Image Enhancement

Three point operations are implemented for image enhancement:

1) Log Transform: The logarithmic transformation enhances dark regions:

$$s = c \cdot \log(1 + r) \tag{2}$$

where r is the input pixel value normalized to [0,1], s is the output, and c is a constant.

2) Gamma Transform: Power-law transformation adjusts overall brightness:

$$s = r^{\gamma} \tag{3}$$

Different γ values produce different effects:

- $\gamma < 1$: Brightens the image
- $\gamma > 1$: Darkens the image
- $\gamma = 2.2$: Standard monitor correction

3) Image Negative: The negative transformation inverts pixel values:

$$s = L - 1 - r \tag{4}$$

where L is the maximum gray level (256 for 8-bit images).

C. Image Resampling

Two interpolation methods are implemented for image resizing:

1) Nearest Neighbor Interpolation: The simplest method assigns the nearest pixel value:

$$I'(x', y') = I(\lfloor x \rfloor, \lfloor y \rfloor) \tag{5}$$

where (x, y) are the corresponding source coordinates.

2) Bilinear Interpolation: This method uses weighted average of four neighboring pixels:

$$\begin{aligned}
I'(x', y') = {} & (1 - dx)(1 - dy)I(x_1, y_1) \\
& + dx(1 - dy)I(x_2, y_1) \\
& + (1 - dx)dyI(x_1, y_2) \\
& + dx \cdot dy \cdot I(x_2, y_2)
\end{aligned} \tag{6}$$

where $dx = x - x_1$ and $dy = y - y_1$ are the fractional parts.

III. Implementation Details

A. Software Architecture

The implementation uses Python with NumPy for numerical operations, PIL for image I/O, and Matplotlib for visualization. The complete implementation is available as open source [6]. The system has been designed with two architectural approaches:

1) Original Monolithic Design: The initial implementation (`hw1_image_processing.py`) contains all functionality in a single ImageProcessor class, providing a straightforward approach for the assignment requirements.

2) Modular Architecture (Extended Implementation): To facilitate future expansion and integration with subsequent assignments, a modular architecture has been developed:

- **Core Module** (`core/`): Contains the ImageProcessor class with all image processing algorithms
- **Utils Module** (`utils/`): Provides reusable utilities
  - `FileIO`: Unified file reading/writing interface supporting multiple formats
  - `Visualizer`: Enhanced visualization tools with histogram and comparison functions
- **Main Interface** (`main.py`): Command-line interface with argument parsing for flexible execution

This modular design offers several advantages:

1) **Reusability**: Core functions can be imported and used in future assignments
2) **Maintainability**: Separation of concerns makes code easier to understand and modify
3) **Extensibility**: New processing modules can be added without modifying existing code
4) **Testing**: Individual modules can be tested independently

B. Key Implementation Components

1) RAW Image Reading: The critical component for reading RAW format images involves binary data handling:

```
def read_raw_image(self, filename,
                   width=512, height=512):
    with open(file_path, 'rb') as f:
        raw_data = f.read()
    img = np.frombuffer(raw_data,
                        dtype=np.uint8)
    img = img.reshape((height, width))
    return img
```

Key insight: RAW files store pixels in row-major order without headers.

2) Bilinear Interpolation Core: The heart of smooth image resizing:

```
# Calculate fractional parts
dx = src_x - x1
dy = src_y - y1

# Weighted average of 4 neighbors
value = (1-dx)*(1-dy)*img[y1,x1] + \
        dx*(1-dy)*img[y1,x2] + \
        (1-dx)*dy*img[y2,x1] + \
        dx*dy*img[y2,x2]
```

This ensures smooth transitions between pixels.

C. Implementation Challenges and Solutions

1) Challenge 1: RAW Format Ambiguity: **Problem:** RAW files lack headers specifying dimensions or bit depth. **Solution:** Assumed standard 512×512 8-bit grayscale format based on file size (262,144 bytes). Implemented parameterized dimensions for flexibility.

2) Challenge 2: Numerical Instability in Log Transform: **Problem:** $\log(0)$ produces undefined results for black pixels. **Solution:** Added 1 to normalized values before applying logarithm:

```
img_normalized = img / 255.0
log_img = c * np.log(1 + img_normalized)
```

3) Challenge 3: Boundary Artifacts in Interpolation: **Problem:** Index out-of-bounds errors when interpolating near edges. **Solution:** Implemented boundary clamping:

```
x2 = min(x1 + 1, old_width - 1)
y2 = min(y1 + 1, old_height - 1)
```

4) Challenge 4: Aspect Ratio Distortion: **Problem:** Non-uniform scaling produced unexpected stretching. **Solution:** Separate scaling factors for width and height:

```
scale_x = (old_width - 1) / (new_width - 1)
scale_y = (old_height - 1) / (new_height - 1)
```

5) Challenge 5: Performance vs. Clarity: **Problem:** Nested loops in Python are significantly slower than vectorized operations. **Solution:** Chose readable implementation for educational clarity, documented optimization opportunities for future work.

IV. EXPERIMENTAL RESULTS

A. Test Images

Six test images were used:

- RAW format: lena, goldhill, peppers (512×512, 8-bit grayscale)
- BMP format: boat, baboon, F16 (512×512, converted to grayscale)

B. Image Reading Results

All images were successfully loaded and displayed. The center 10×10 pixel values for each image are shown below for verification as required.

TABLE I
CENTER 10×10 PIXEL VALUES FOR LENA.RAW

Lena.raw									
195	195	195	192	169	135	133	137	138	148
196	196	189	157	124	128	135	144	145	145
196	187	149	123	127	131	135	142	142	137
180	135	116	117	124	130	131	132	137	133
124	102	115	116	120	126	124	120	114	111
100	102	114	114	114	118	120	117	112	92
101	100	113	106	98	89	90	104	101	84
105	96	99	94	92	78	75	79	83	76
92	90	90	84	79	75	71	73	72	71
85	79	76	77	76	71	73	73	69	77

C. Enhancement Results

Figure 1 shows the enhancement results for one test image. Key observations:

TABLE II
CENTER 10×10 PIXEL VALUES FOR ALL TEST IMAGES

Image	Center 10×10 Pixel Values (First Row)
Goldhill.raw	56, 60, 69, 64, 54, 74, 80, 78, 78, 121
Peppers.raw	63, 43, 45, 60, 71, 61, 58, 25, 4, 20
Boat.bmp	88, 74, 105, 105, 48, 57, 62, 75, 140, 142
Baboon.bmp	152, 151, 154, 167, 155, 142, 139, 136, 139, 143
F16.bmp	105, 110, 108, 111, 118, 130, 126, 123, 115, 107

Fig. 1. Image enhancement results on Lena image

1) Log Transform:

- Effectively brightens dark regions
- Compresses dynamic range
- Useful for images with low contrast in dark areas

2) Gamma Transform:

- $\gamma = 0.5$: Significantly brightens the image
- $\gamma = 1.5$: Slightly darkens while preserving details
- $\gamma = 2.2$: Standard display correction

3) Negative Transform:

- Complete inversion of gray levels
- Useful for enhancing white or gray details in dark regions

D. Resampling Results

Table III summarizes the resampling experiments:

TABLE III
RESAMPLING TEST CASES

Test Case	Source Size	Target Size
Downsampling 1	512×512	128×128
Downsampling 2	512×512	32×32
Upsampling 1	32×32	512×512
Non-uniform 1	512×512	1024×512
Non-uniform 2	128×128	256×512

1) Downsampling Results:

- Both methods preserve main features
- Bilinear produces smoother results
- Nearest neighbor shows more aliasing artifacts

2) Upsampling Results:

- Nearest neighbor produces blocky artifacts
- Bilinear creates smoother but blurrier results
- Information loss is irreversible

Figure 2 shows a comparison of the two methods:

Fig. 2. Comparison of interpolation methods (32×32 to 512×512)

V. DISCUSSION

A. Lessons Learned

1) Understanding Image Formats: Working with RAW format taught us the importance of metadata. Unlike modern formats (JPEG, PNG) that include headers, RAW files require external knowledge of image dimensions. This experience highlighted why standardized formats with self-describing headers became prevalent.

2) Trade-offs in Algorithm Design: The implementation revealed several key trade-offs:

- **Quality vs. Speed**: Bilinear interpolation produces better results but requires 4× more computations
- **Memory vs. Computation**: Pre-computing scaling factors reduces redundant calculations
- **Readability vs. Performance**: Educational code prioritizes understanding over optimization

3) Practical Insights:

1) Gamma correction with $\gamma = 2.2$ is not arbitrary—it compensates for typical monitor characteristics
2) Log transform is particularly effective for images with wide dynamic range
3) Upsampling cannot recover lost information; the choice of interpolation only affects perceived quality

B. Performance Analysis

Bilinear interpolation requires approximately 4× more computations than nearest neighbor but produces significantly better visual quality. The trade-off between quality and speed depends on the application requirements. Our testing showed:

- Nearest neighbor: 0.05 seconds for 512×512 resize
- Bilinear: 0.20 seconds for 512×512 resize
- Quality difference most noticeable in upsampling scenarios

C. Limitations

- Current implementation uses simple loops, which could be optimized with vectorization
- Only grayscale images are supported
- More advanced interpolation methods (bicubic, Lanczos) could improve quality

D. Potential Improvements and Achievements

1) Completed Improvements: The following enhancements have been implemented in the extended version:

- **Modular Architecture**: Successfully refactored into reusable modules
- **Enhanced CLI**: Added command-line arguments for flexible operation
- **Unified I/O Interface**: Created FileIO class for consistent file handling
- **Extended Visualization**: Implemented additional visualization utilities

2) Future Improvements:

- Implement vectorized operations for better performance
- Add support for color images
- Include more interpolation methods (bicubic, Lanczos)
- Implement adaptive enhancement based on image statistics
- Develop a graphical user interface (GUI) for interactive processing

VI. CONCLUSION

This assignment successfully implemented fundamental image processing operations including reading, enhancement, and resampling. The results demonstrate:

1) Correct reading of both RAW and standard image formats
2) Effective enhancement through point operations, each suitable for different scenarios
3) Clear differences between interpolation methods in resampling tasks
4) Successful modular architecture design for future extensibility

The implementation provides a solid foundation for understanding digital image processing concepts. The visual and quantitative results confirm the theoretical expectations for each technique. Furthermore, the modular architecture ensures that these implementations can be readily integrated with future assignments, creating a comprehensive image processing system.

The extended implementation with its modular design demonstrates software engineering best practices, including separation of concerns, reusability, and maintainability. This approach not only fulfills the current assignment requirements but also establishes a framework for building more complex image processing applications.

Future work will leverage this modular foundation to incorporate more sophisticated algorithms such as filtering operations, edge detection, morphological processing, and advanced enhancement techniques.

ACKNOWLEDGMENT

REFERENCES

[1] R. C. Gonzalez and R. E. Woods, *Digital Image Processing*, 4th ed. Pearson, 2018.

[2] Wikipedia contributors, "Bilinear interpolation," *Wikipedia, The Free Encyclopedia*. [Online]. Available: https://en.wikipedia.org/wiki/Bilinear_interpolation

[3] Wikipedia contributors, "Nearest-neighbor interpolation," *Wikipedia, The Free Encyclopedia*. [Online]. Available: https://en.wikipedia.org/wiki/Nearest-neighbor_interpolation

[4] NumPy Developers, "NumPy: The fundamental package for scientific computing with Python," 2023. [Online]. Available: https://numpy.org/

[5] F. Lundh and Contributors, "Python Imaging Library (PIL)," 2023. [Online]. Available: https://python-pillow.org/

[6] I. Tsai, "MMIP-HW1: Digital Image Processing Implementation," 2025. [Online]. Available: https://github.com/IaTsai/MMIP-HW1.git